

Homework 8

Author: Anthony Martinez | amm180005 \ Date: 11/19/21

1. Read the Auto data (5 points)

- use pandas to read the data (available on Piazza)
- print the first few rows
- print the dimensions of the data

```
In [172]: import pandas as pd

# a. use pandas to read the data
df = pd.read_csv('Auto.csv')

# b. print the first few rows
df.head()

# c. print the dimensions of the data
print('Dimensions of the data frame: ', df.shape)
```

Dimensions of the data frame: (392, 9)

2. Some data exploration with code (10 points)

- use describe() on the mpg, weight, and year columns
- write comments indicating the range and average of each column

Commentary:

mpg: range = (9.0, 46.0) | average = 23.445918

weight: range = (1613.0, 5140.0) | average = 2977.584184

year: range = (70.0, 82.0) | average = 76.010256

```
In [173]: # a. use describe() on the mpg, weight, and year columns  
df.loc[:, ['mpg', 'weight', 'year']].describe()
```

Out[173]:

| | mpg | weight | year |
|-------|------------|-------------|------------|
| count | 392.000000 | 392.000000 | 390.000000 |
| mean | 23.445918 | 2977.584184 | 76.010256 |
| std | 7.805007 | 849.402560 | 3.668093 |
| min | 9.000000 | 1613.000000 | 70.000000 |
| 25% | 17.000000 | 2225.250000 | 73.000000 |
| 50% | 22.750000 | 2803.500000 | 76.000000 |
| 75% | 29.000000 | 3614.750000 | 79.000000 |
| max | 46.600000 | 5140.000000 | 82.000000 |

3. Explore data types (10 points)

- check the data types of all columns
- change the cylinders column to categorical (use cat.codes)
- change the origin column to categorical (don't use cat.codes)
- verify the changes with the dtypes attribute

```
In [174]: # a. check the data types of all columns
print('Before column conversions:\n')
print(df.dtypes)

# b. change the cylinders column to categorical (use cat.codes)
df.cylinders = df.cylinders.astype('category').cat.codes
# using cat.codes will mean data type will show up as int8 rather than category

# c. change the origin column to categorical (don't use cat.codes)
df.origin = df.origin.astype('category')

# d. verify the changes with the dtypes attribute
print('\nAfter column conversions:\n')
print(df.dtypes)
```

Before column conversions:

| | |
|--------------|---------|
| mpg | float64 |
| cylinders | int64 |
| displacement | float64 |
| horsepower | int64 |
| weight | int64 |
| acceleration | float64 |
| year | float64 |
| origin | int64 |
| name | object |

dtype: object

After column conversions:

| | |
|--------------|----------|
| mpg | float64 |
| cylinders | int8 |
| displacement | float64 |
| horsepower | int64 |
| weight | int64 |
| acceleration | float64 |
| year | float64 |
| origin | category |
| name | object |

dtype: object

4. Deal with NAs (5 points)

- delete rows with NAs
- print the new dimensions

```
In [175]: # check for NAs
df.isnull().sum()

print('Acceleration and year are the only columns with NAs')

# a. delete rows with NAs
df = df.dropna()

# b. print the new dimensions / new dimensions indicate that the rows were deleted
print('\nDimensions of data frame:', df.shape)
```

Acceleration and year are the only columns with NAs

Dimensions of data frame: (389, 9)

5. Modify columns (10 points)

- a. make a new column, mpg_high, which is categorical:
 - i. the column == 1 if mpg > average mpg, else == 0
- b. delete the mpg and name columns
- c. print the first few rows of the modified data frame

```
In [176]: import numpy as np

# a. make a new column, mpg_high, which is categorical: the column == 1 if mpg
# > average mpg, else == 0
average_mpg = np.mean(df.mpg)
df['mpg_high'] = np.where(df['mpg'] > average_mpg , 1, 0)

# b. delete the mpg and name columns
df1 = df.copy()
df1 = df1.drop(columns=['mpg', 'name'])

# c. print the first few rows of the modified data frame
df1.head()
```

Out[176]:

| | cylinders | displacement | horsepower | weight | acceleration | year | origin | mpg_high |
|---|-----------|--------------|------------|--------|--------------|------|--------|----------|
| 0 | 4 | 307.0 | 130 | 3504 | 12.0 | 70.0 | 1 | 0 |
| 1 | 4 | 350.0 | 165 | 3693 | 11.5 | 70.0 | 1 | 0 |
| 2 | 4 | 318.0 | 150 | 3436 | 11.0 | 70.0 | 1 | 0 |
| 3 | 4 | 304.0 | 150 | 3433 | 12.0 | 70.0 | 1 | 0 |
| 6 | 4 | 454.0 | 220 | 4354 | 9.0 | 70.0 | 1 | 0 |

6. Data exploration with graphs (15 points)

- seaborn catplot on the mpg_high column
- seaborn relplot with horsepower on the x axis, weight on the y axis, setting hue or style to mpg_high
- seaborn boxplot with mpg_high on the x axis and weight on the y axis
- for each graph, write a comment indicating one thing you learned about the data from the graph

Commentary:

Seaborn Catplot of mpg_high: From this graph we can see that there are more instances with a mpg below the average mpg than instances that have an mpg higher than the average

Seaborn Relplot of horsepower and weight: From this graph we learned that automobiles that have mpg score higher than the average, tend to have a lower horsepower score. In other words, automobiles with a mpg score lower than the average, and that weight more, tend to have a higher horsepower score.



Seaborn Boxplot of mpg_high and weight: From this graph we learned that automobiles with a higher weight tend to have lower mpg scores. Rather, lighter automobiles tend to have a higher mpg score.

```
In [177]: import seaborn as sb

# a. seaborn catplot on the mpg_high column
sb.catplot(x='mpg_high', kind='count', data=df1)


# b. seaborn relplot with horsepower on the x axis, weight on the y axis, setting hue or style to mpg_high
sb.relplot(x='horsepower', y='weight', data=df1, hue=df1.mpg_high)
```

Out[177]: <seaborn.axisgrid.FacetGrid at 0x17e4f6d9888>



In [178]: *# c. seaborn boxplot with mpg_high on the x axis and weight on the y axis*
`sb.boxplot(x='mpg_high', y='weight', data=df1)`

Out[178]: `<matplotlib.axes._subplots.AxesSubplot at 0x17e4df58588>`



7. Train/test split (5 points)

- 80/20
- use seed 1234 so we all get the same results
- train /test X data frames consists of all remaining columns except mpg_high
- print the dimensions of train and test

In [179]: `from sklearn.model_selection import train_test_split`

c. train /test X data frames consists of all remaining columns except mpg_high
`X = df1.loc[:, ['cylinders', 'displacement', 'horsepower', 'weight', 'acceleration', 'year', 'origin']]`
`y = df1.mpg_high`

a. 80/20, b. use seed 1234 so we all get the same results
`X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=1234)`

d. print the dimensions of train and test
`print('train size:', X_train.shape)`
`print('test size:', X_test.shape)`

train size: (311, 7)
test size: (78, 7)

8. Logistic Regression (10 points)

- train a logistic regression model using solver lbfgs
- test and evaluate
- print metrics using the classification report

```
In [180]: from sklearn.linear_model import LogisticRegression
import warnings
warnings.filterwarnings('ignore')

glm1 = LogisticRegression(solver='lbfgs')

# a. train a logistic regression model using solver lbfgs
glm1.fit(X_train, y_train)
glm1.score(X_train, y_train)

# b. test and evaluate
pred = glm1.predict(X_test)

# c. print metrics using the classification report
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix
print('accuracy:', accuracy_score(y_test, pred))
print('precision:', precision_score(y_test, pred))
print('recall score:', recall_score(y_test, pred))
print('f1 score', f1_score(y_test, pred))
print('Confusion Matrix:')
confusion_matrix(y_test, pred)

accuracy: 0.8589743589743589
precision: 0.7297297297297297
recall score: 0.9642857142857143
f1 score 0.8307692307692307
Confusion Matrix:
```

```
Out[180]: array([[40, 10],
                [ 1, 27]], dtype=int64)
```

9. Decision Tree (10 points)

- train a decision tree
- test and evaluate
- print the classification report metrics
- plot the tree (optional, see: <https://scikit-learn.org/stable/modules/tree.html>)


```
In [181]: from sklearn.tree import DecisionTreeClassifier
          from sklearn import tree

          # a. train a decision tree
          dt = DecisionTreeClassifier()
          dt.fit(X_train, y_train)

          # b. test and evaluate
          pred2 = dt.predict(X_test)

          # c. print the classification report metrics
          print('accuracy score: ', accuracy_score(y_test, pred2))
          print('precision score: ', precision_score(y_test, pred2))
          print('recall score: ', recall_score(y_test, pred2))
          print('f1 score: ', f1_score(y_test, pred2))
          print('Confusion matrix:')
          print(confusion_matrix(y_test, pred2))

          # d. plot the tree
          tree.plot_tree(dt)
```

```
accuracy score: 0.9230769230769231
precision score: 0.8666666666666667
recall score: 0.9285714285714286
f1 score: 0.896551724137931
Confusion matrix:
[[46  4]
 [ 2 26]]
```

```


Out[181]: [Text(215.40441176470588, 205.35999999999999, 'X[0] <= 2.5\ngini = 0.5\nsamples = 311\nvalue = [153, 158]'),
Text(145.24411764705883, 181.2, 'X[2] <= 101.0\ngini = 0.239\nsamples = 173\nvalue = [24, 149]'),
Text(93.54705882352941, 157.04, 'X[5] <= 75.5\ngini = 0.179\nsamples = 161\nvalue = [16, 145]'),
Text(49.23529411764706, 132.88, 'X[1] <= 119.5\ngini = 0.362\nsamples = 59\nvalue = [14, 45]'),
Text(19.694117647058825, 108.72, 'X[4] <= 13.75\ngini = 0.159\nsamples = 46\nvalue = [4, 42]'),
Text(9.847058823529412, 84.56, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(29.541176470588237, 84.56, 'X[3] <= 2683.0\ngini = 0.087\nsamples = 44\nvalue = [2, 42]'),
Text(19.694117647058825, 60.400000000000006, 'X[3] <= 2377.0\ngini = 0.045\nsamples = 43\nvalue = [1, 42]'),
Text(9.847058823529412, 36.240000000000001, 'gini = 0.0\nsamples = 38\nvalue = [0, 38]'),
Text(29.541176470588237, 36.240000000000001, 'X[3] <= 2385.0\ngini = 0.32\nsamples = 5\nvalue = [1, 4]'),
Text(19.694117647058825, 12.079999999999984, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(39.38823529411765, 12.079999999999984, 'gini = 0.0\nsamples = 4\nvalue = [0, 4]'),
Text(39.38823529411765, 60.400000000000006, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(78.7764705882353, 108.72, 'X[4] <= 17.75\ngini = 0.355\nsamples = 13\nvalue = [10, 3]'),
Text(68.92941176470589, 84.56, 'X[2] <= 81.5\ngini = 0.469\nsamples = 8\nvalue = [5, 3]'),
Text(59.082352941176474, 60.400000000000006, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]'),
Text(78.7764705882353, 60.400000000000006, 'X[1] <= 131.0\ngini = 0.278\nsamples = 6\nvalue = [5, 1]'),
Text(68.92941176470589, 36.240000000000001, 'gini = 0.0\nsamples = 4\nvalue = [4, 0]'),
Text(88.62352941176471, 36.240000000000001, 'X[2] <= 86.5\ngini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(78.7764705882353, 12.079999999999984, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(98.47058823529412, 12.079999999999984, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(88.62352941176471, 84.56, 'gini = 0.0\nsamples = 5\nvalue = [5, 0]'),
Text(137.85882352941178, 132.88, 'X[3] <= 3250.0\ngini = 0.038\nsamples = 102\nvalue = [2, 100]'),
Text(118.16470588235295, 108.72, 'X[3] <= 2880.0\ngini = 0.02\nsamples = 100\nvalue = [1, 99]'),
Text(108.31764705882354, 84.56, 'gini = 0.0\nsamples = 94\nvalue = [0, 94]'),
Text(128.01176470588237, 84.56, 'X[3] <= 2920.0\ngini = 0.278\nsamples = 6\nvalue = [1, 5]'),
Text(118.16470588235295, 60.400000000000006, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(137.85882352941178, 60.400000000000006, 'gini = 0.0\nsamples = 5\nvalue = [0, 5]'),
Text(157.5529411764706, 108.72, 'X[1] <= 151.5\ngini = 0.5\nsamples = 2\nvalue = [1, 1]'),
Text(147.7058823529412, 84.56, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),

```

```

Text(167.4, 84.56, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(196.94117647058823, 157.04, 'X[4] <= 14.45\ngini = 0.444\nsamples = 12\nvalue = [8, 4]'),
Text(187.09411764705882, 132.88, 'X[5] <= 76.0\ngini = 0.444\nsamples = 6\nvalue = [2, 4]'),
Text(177.24705882352941, 108.72, 'gini = 0.0\nsamples = 3\nvalue = [0, 3]'),
Text(196.94117647058823, 108.72, 'X[3] <= 2760.0\ngini = 0.444\nsamples = 3\nvalue = [2, 1]'),
Text(187.09411764705882, 84.56, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(206.78823529411767, 84.56, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(206.78823529411767, 132.88, 'gini = 0.0\nsamples = 6\nvalue = [6, 0]'),
Text(285.56470588235294, 181.2, 'X[5] <= 79.5\ngini = 0.122\nsamples = 138\nvalue = [129, 9]'),
Text(265.8705882352941, 157.04, 'X[4] <= 21.6\ngini = 0.045\nsamples = 129\nvalue = [126, 3]'),
Text(256.02352941176474, 132.88, 'X[3] <= 2737.0\ngini = 0.031\nsamples = 128\nvalue = [126, 2]'),
Text(236.3294117647059, 108.72, 'X[3] <= 2674.0\ngini = 0.444\nsamples = 3\nvalue = [2, 1]'),
Text(226.4823529411765, 84.56, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(246.1764705882353, 84.56, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(275.71764705882356, 108.72, 'X[2] <= 83.0\ngini = 0.016\nsamples = 125\nvalue = [124, 1]'),
Text(265.8705882352941, 84.56, 'X[2] <= 79.5\ngini = 0.375\nsamples = 4\nvalue = [3, 1]'),
Text(256.02352941176474, 60.400000000000006, 'gini = 0.0\nsamples = 3\nvalue = [3, 0]'),
Text(275.71764705882356, 60.400000000000006, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(285.56470588235294, 84.56, 'gini = 0.0\nsamples = 121\nvalue = [121, 0]'),
Text(275.71764705882356, 132.88, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(305.25882352941176, 157.04, 'X[1] <= 196.5\ngini = 0.444\nsamples = 9\nvalue = [3, 6]'),
Text(295.4117647058824, 132.88, 'gini = 0.0\nsamples = 4\nvalue = [0, 4]'),
Text(315.1058823529412, 132.88, 'X[1] <= 247.0\ngini = 0.48\nsamples = 5\nvalue = [3, 2]'),
Text(305.25882352941176, 108.72, 'gini = 0.0\nsamples = 3\nvalue = [3, 0]'),
Text(324.95294117647063, 108.72, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]')

```



10. Analysis (20 points)

- a. which algorithm performed better?
- b. compare accuracy, recall and precision metrics by class
- c. give your analysis of why the better-performing algorithm might have outperformed the other

Commentary:

The decision tree algorithm performed better than the logistic regression algorithm. It had a higher accuracy score (92 vs 85), a higher precision score (86 vs 72), but a slightly smaller recall score (92.8 vs 96.6). The reason for this outcome is most likely due to the nature of the data. Logistic Regression models will tend to perform better on data that is linearly separable. On the other hand, Decision Trees are non-linear classifiers which means they do not require the data to be linearly separable, allowing them to outperform logistic regression models when the data is not linearly separable.