

# Homework 9 (Optional)

author: Anthony Martinez \ CS 4375

## Load The Data

In [140...

```
import pandas as pd
df0 = pd.read_csv('data/blackfriday.csv')
```

## Data Cleaning

Link to Dataset: <https://www.kaggle.com/sdolezel/black-friday?select=train.csv> Data Cleaning Operations:

1. Count NAs, if any
2. Remove NAs
3. Convert qualitative data into factors

In [141...

```
# 1. Count NAs
df0.isnull().sum()
# Product_Category_2 and Product_Category_3 are the only columns with NA

# 2. Remove NAs
df = df0.drop(columns=['Product_Category_2', 'Product_Category_3'])

# confirm removal of columns
df.isnull().sum()

# 3. Convert qualitative data into factors
df.dtypes

'''
The following columns will be converted to category type using cat codes
- Gender
- Age
- Occupation
- Stay_In_Current_City_Years
- City_Category
- Marital_Status
- Product_Category
'''

df1 = df.copy()
df1.Gender = df1.Gender.astype('category').cat.codes
df1.Age = df1.Age.astype('category').cat.codes
df1.Occupation = df1.Occupation.astype('category').cat.codes
df1.City_Category = df1.City_Category.astype('category').cat.codes
df1.Stay_In_Current_City_Years = df1.Stay_In_Current_City_Years.astype('category').cat.codes
df1.Marital_Status = df1.Marital_Status.astype('category').cat.codes
df1.Product_Category_1 = df1.Product_Category_1.astype('category').cat.codes
```

```
# Checking that all columns were set to category type
df1.dtypes
```

```
Out[141... User_ID          int64
Product_ID       object
Gender           int8
Age             int8
Occupation       int8
City_Category    int8
Stay_In_Current_City_Years  int8
Marital_Status   int8
Product_Category_1  int8
Purchase         int64
dtype: object
```

## Data Exploration

- use at least 5 R functions for data exploration
- create at least 3 informative R graphs for data exploration

```
In [142... ...
5 R functions
1. head
2. info
3. info
4. columns
5. mean(purchase)
...

# same as R's head() function
print(df1.head())

# describe() is similar to the summary() in R
print('\n',df1.describe())

# info() is similar to R's str()
print('\n', df1.info())

# .columns is similar to colnames(df) in R
print('\n', df1.columns)

print('Mean of Purchase column (target column)', df1["Purchase"].mean())
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	\
0	1000001	P00069042	0	0	10	0	
1	1000001	P00248942	0	0	10	0	
2	1000001	P00087842	0	0	10	0	
3	1000001	P00085442	0	0	10	0	
4	1000002	P00285442	1	6	16	2	

	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	Purchase
0	2	0	2	8370
1	2	0	0	15200
2	2	0	11	1422
3	2	0	11	1057
4	4	0	7	7969

	User_ID	Gender	Age	Occupation	\
--	---------	--------	-----	------------	---

count	5.500680e+05	550068.000000	550068.000000	550068.000000
mean	1.003029e+06	0.753105	2.496430	8.076707
std	1.727592e+03	0.431205	1.353632	6.522660
min	1.000001e+06	0.000000	0.000000	0.000000
25%	1.001516e+06	1.000000	2.000000	2.000000
50%	1.003077e+06	1.000000	2.000000	7.000000
75%	1.004478e+06	1.000000	3.000000	14.000000
max	1.006040e+06	1.000000	6.000000	20.000000

	City_Category	Stay_In_Current_City_Years	Marital_Status	\
count	550068.000000	550068.000000	550068.000000	
mean	1.042640	1.858418	0.409653	
std	0.760211	1.289443	0.491770	
min	0.000000	0.000000	0.000000	
25%	0.000000	1.000000	0.000000	
50%	1.000000	2.000000	0.000000	
75%	2.000000	3.000000	1.000000	
max	2.000000	4.000000	1.000000	

	Product_Category_1	Purchase
count	550068.000000	550068.000000
mean	4.404270	9263.968713
std	3.936211	5023.065394
min	0.000000	12.000000
25%	0.000000	5823.000000
50%	4.000000	8047.000000
75%	7.000000	12054.000000
max	19.000000	23961.000000

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 550068 entries, 0 to 550067

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
0	User_ID	550068 non-null	int64
1	Product_ID	550068 non-null	object
2	Gender	550068 non-null	int8
3	Age	550068 non-null	int8
4	Occupation	550068 non-null	int8
5	City_Category	550068 non-null	int8
6	Stay_In_Current_City_Years	550068 non-null	int8
7	Marital_Status	550068 non-null	int8
8	Product_Category_1	550068 non-null	int8
9	Purchase	550068 non-null	int64

dtypes: int64(2), int8(7), object(1)

memory usage: 16.3+ MB

None

```
Index(['User_ID', 'Product_ID', 'Gender', 'Age', 'Occupation', 'City_Category',
      'Stay_In_Current_City_Years', 'Marital_Status', 'Product_Category_1',
      'Purchase'],
      dtype='object')
```

Mean of Purchase column (target column) 9263.968712959126


## Data Exploration: Graphs

In [143...

```
import seaborn as sb
```

```
# 1. seaborn catplot on the purchase column
sb.catplot(x='Purchase', kind='count', data=df1)
```


Out[143... <seaborn.axisgrid.FacetGrid at 0x7fb456e289d0>



In [144... `import seaborn as sb`

```
# 2. seaborn boxplot with Occupation on the x axis and Purchase on the y axis
sb.boxplot(x='Occupation', y='Purchase', data=df1)
```


Out[144... <AxesSubplot:xlabel='Occupation', ylabel='Purchase'>



In [162... `import seaborn as sb`

```
# 3. seaborn boxplot with Age on the x axis and Purchase on the y axis
sb.boxplot(x='Age', y='Purchase', data=df1)
```

Out[162... <AxesSubplot:xlabel='Age', ylabel='Purchase'>



## Train/Test Split (80/20)

In [145...

```
from sklearn.model_selection import train_test_split

# Predicting Purchase from Age, Occupation, City_Category, and Product_Category_1
X = df1.loc[:, ['Age', 'Occupation', 'City_Category', 'Product_Category_1']]
y = df1.Purchase

# 80/20 split, using seed 1234 for repeatability
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_

# printing sizes of train and test
print('train size:', X_train.shape)
print('test size:', X_test.shape)
```

train size: (440054, 4)

test size: (110014, 4)

## Algorithm 1: Linear Regression

In [153...

```
from sklearn.linear_model import LinearRegression

linreg = LinearRegression()
linreg.fit(X_train, y_train)

# make predictions
y_pred = linreg.predict(X_test)

# evaluation
from sklearn.metrics import mean_squared_error, r2_score
print('mse=', mean_squared_error(y_test, y_pred))
print('correlation=', r2_score(y_test, y_pred))
```

mse= 22126347.338500306

correlation= 0.12262488554739837

## Algorithm 2: Decision Tree

In [158...

```
from sklearn.tree import DecisionTreeRegressor
```

```

from sklearn import tree

regressor = DecisionTreeRegressor()

regressor.fit(X_train, y_train)
pred_dt = regressor.predict(X_test)

# evaluation
from sklearn.metrics import mean_squared_error, r2_score
print('mse=', mean_squared_error(y_test, pred_dt))
print('correlation=', r2_score(y_test, pred_dt))

```

```

mse= 8999419.148505524
correlation= 0.6431464134304717

```

## Algorithm 3: KNN Regression

In [160...

```

# train the algorithm
from sklearn.neighbors import KNeighborsRegressor
regressor2 = KNeighborsRegressor()
regressor2.fit(X_train, y_train)

# make predictions
pred_knn = regressor2.predict(X_test)

# evaluation
from sklearn.metrics import mean_squared_error, r2_score
print('mse=', mean_squared_error(y_test, y_pred))
print('correlation=', r2_score(y_test, y_pred))

```

```

mse= 10476668.805379316
correlation= 0.5845690953152732

```

## Analysis:

The linear regression algorithm got a mse of 22126347.34 and a correlation= 0.1226 The Decision tree regression algorithm got a mse of 8999419.15 and a correlation of 0.64 The kNN regression algorithm got a mse mse of 10476668.81 and a correlation of 0.585 Rank:

1. Decision Tree Regression
2. kNN Regression
3. Linear Regression (multiple)

Decision Tree Regression did the best out of all three algorithms. KNN preformed the 2nd best. This is probably because the algorithm uses feature similarity to make predictions. The reason linear regression probably did poorly is because of its high bias. The predictor data varies wildly. Linear Regression will find a line, even if it makes no sense.

## Machine Learning in R v.s Python

Coming into this course I thought I would enjoy using Python more than R. I was a little upset when I found out this course is taught in R and not Python. Python is my favorite language to build projects and it still is. However, when it comes to Machine Learning, data cleaning, and

data exploration R is more intuitive. Another important factor is R-Studio. R-Studio is a great IDE for ML and for working with data. I tried to find an IDE similar to R-Studio, but for Python. The best option I came across was DataSpell. The IDE is pretty great but I think it's safe to say that R-Studio is a better option. I think it is also important to note that we did a lot more assignments using R than in comparison to Python, so it's not much of a fair comparison. One thing I have noticed, especially towards the end of this course, is that it seems Machine Learning with Python seems to have more powerful tools for Deep Learning such as Keras and TensorFlow. It also seems that these libraries are in high demand when it comes to the current job posting. It is for this last reason that I have decided to continue my Deep Learning and Machine Learning studies using Python even though my personal opinion is that R/R-Studio is more user friendly than Python.