

R Project - Classification

4375 Machine Learning with Dr. Mazidi

Anthony Martinez | netid: amm180005

10/17/21

Project: Using data from NBA games spanning from 2004-2015 I will use ML classification algorithms to predict the winner.

The data set was downloaded from Kaggle: <https://www.kaggle.com/nathanlauga/nba-games>
(<https://www.kaggle.com/nathanlauga/nba-games>)

Load the data

```
df0<-read.csv("data/games.csv", header=TRUE)
```

Data Exploration

- use at least 5 R functions for data exploration
- create at least 2 informative R graphs for data exploration

I am a little surprised that the mean of points scored for home teams was only 2.8 points higher than the mean of points scored for away teams. I was planning on using home court advantage as a predictor, but maybe it won't make the best predictor.

```
# Use head to get a peak at the first 5 observations in the data
head(df0)
```

GAME_DATE... <chr>	GAME... <int>	GAME_STATUS_... <chr>	HOME_TEA... <int>	VISITOR_TEAM_ID <int>	SEA... <int>	TEAM_I
1 5/26/2021	42000102	Final	1610612755	1610612764	2020	16106
2 5/26/2021	42000132	Final	1610612752	1610612737	2020	16106
3 5/26/2021	42000142	Final	1610612762	1610612763	2020	16106
4 5/25/2021	42000112	Final	1610612751	1610612738	2020	16106
5 5/25/2021	42000152	Final	1610612756	1610612747	2020	16106
6 5/25/2021	42000172	Final	1610612746	1610612742	2020	16106

6 rows | 1-8 of 22 columns

```
# Using summary to get statistics for each column
summary(df0)
```

```
## GAME_DATE_EST      GAME_ID      GAME_STATUS_TEXT      HOME_TEAM_ID
## Length:24677      Min.   :10300001      Length:24677      Min.   :1.611e+09
## Class :character  1st Qu.:20600878      Class :character  1st Qu.:1.611e+09
## Mode  :character  Median :21100897      Mode  :character  Median :1.611e+09
##                               Mean   :21644558      Mean   :1.611e+09
##                               3rd Qu.:21601157      3rd Qu.:1.611e+09
##                               Max.    :52000211      Max.    :1.611e+09
##
## VISITOR_TEAM_ID      SEASON      TEAM_ID_home      PTS_home
## Min.   :1.611e+09      Min.   :2003      Min.   :1.611e+09      Min.   : 36.0
## 1st Qu.:1.611e+09      1st Qu.:2007      1st Qu.:1.611e+09      1st Qu.: 94.0
## Median :1.611e+09      Median :2011      Median :1.611e+09      Median :102.0
## Mean   :1.611e+09      Mean   :2011      Mean   :1.611e+09      Mean   :102.8
## 3rd Qu.:1.611e+09      3rd Qu.:2016      3rd Qu.:1.611e+09      3rd Qu.:111.0
## Max.   :1.611e+09      Max.   :2020      Max.   :1.611e+09      Max.   :168.0
##                               NA's    :99
## FG_PCT_home      FT_PCT_home      FG3_PCT_home      AST_home
## Min.   :0.2500      Min.   :0.1430      Min.   :0.0000      Min.   : 6.00
## 1st Qu.:0.4210      1st Qu.:0.6960      1st Qu.:0.2860      1st Qu.:19.00
## Median :0.4590      Median :0.7650      Median :0.3570      Median :22.00
## Mean   :0.4603      Mean   :0.7591      Mean   :0.3561      Mean   :22.65
## 3rd Qu.:0.5000      3rd Qu.:0.8280      3rd Qu.:0.4290      3rd Qu.:26.00
## Max.   :0.6840      Max.   :1.0000      Max.   :1.0000      Max.   :50.00
## NA's    :99          NA's    :99          NA's    :99          NA's    :99
## REB_home      TEAM_ID_away      PTS_away      FG_PCT_away
## Min.   :15.00      Min.   :1.611e+09      Min.   : 33.00      Min.   :0.244
## 1st Qu.:39.00      1st Qu.:1.611e+09      1st Qu.: 91.00      1st Qu.:0.411
## Median :43.00      Median :1.611e+09      Median : 99.00      Median :0.448
## Mean   :43.27      Mean   :1.611e+09      Mean   : 99.91      Mean   :0.449
## 3rd Qu.:48.00      3rd Qu.:1.611e+09      3rd Qu.:109.00      3rd Qu.:0.487
## Max.   :72.00      Max.   :1.611e+09      Max.   :168.00      Max.   :0.674
## NA's    :99          NA's    :99          NA's    :99          NA's    :99
## FT_PCT_away      FG3_PCT_away      AST_away      REB_away
## Min.   :0.1430      Min.   :0.0000      Min.   : 4.0      Min.   :19.00
## 1st Qu.:0.6920      1st Qu.:0.2780      1st Qu.:18.0      1st Qu.:38.00
## Median :0.7620      Median :0.3500      Median :21.0      Median :42.00
## Mean   :0.7574      Mean   :0.3494      Mean   :21.3      Mean   :41.97
## 3rd Qu.:0.8280      3rd Qu.:0.4210      3rd Qu.:25.0      3rd Qu.:46.00
## Max.   :1.0000      Max.   :1.0000      Max.   :46.0      Max.   :81.00
## NA's    :99          NA's    :99          NA's    :99          NA's    :99
## HOME_TEAM_WINS
## Min.   :0.0000
## 1st Qu.:0.0000
## Median :1.0000
## Mean   :0.5891
## 3rd Qu.:1.0000
## Max.   :1.0000
##
```

```
# Using names to get the names of the columns in the data set
names(df0)
```

```
## [1] "GAME_DATE_EST"      "GAME_ID"           "GAME_STATUS_TEXT"  "HOME_TEAM_ID"
## [5] "VISITOR_TEAM_ID"    "SEASON"            "TEAM_ID_home"      "PTS_home"
## [9] "FG_PCT_home"        "FT_PCT_home"       "FG3_PCT_home"      "AST_home"
## [13] "REB_home"           "TEAM_ID_away"      "PTS_away"          "FG_PCT_away"
## [17] "FT_PCT_away"        "FG3_PCT_away"      "AST_away"          "REB_away"
## [21] "HOME_TEAM_WINS"
```

```
# just out of curiosity I wanted to see if there was any difference between names() and colnames
() functions
colnames(df0)
```

```
## [1] "GAME_DATE_EST"      "GAME_ID"           "GAME_STATUS_TEXT"  "HOME_TEAM_ID"
## [5] "VISITOR_TEAM_ID"    "SEASON"            "TEAM_ID_home"      "PTS_home"
## [9] "FG_PCT_home"        "FT_PCT_home"       "FG3_PCT_home"      "AST_home"
## [13] "REB_home"           "TEAM_ID_away"      "PTS_away"          "FG_PCT_away"
## [17] "FT_PCT_away"        "FG3_PCT_away"      "AST_away"          "REB_away"
## [21] "HOME_TEAM_WINS"
```

```
# using str() to get row/column counts and info on each column
str(df0)
```

```
## 'data.frame': 24677 obs. of 21 variables:
## $ GAME_DATE_EST : chr "5/26/2021" "5/26/2021" "5/26/2021" "5/25/2021" ...
## $ GAME_ID : int 42000102 42000132 42000142 42000112 42000152 42000172 42000122 4200
0162 42000101 42000151 ...
## $ GAME_STATUS_TEXT: chr "Final" "Final" "Final" "Final" ...
## $ HOME_TEAM_ID : int 1610612755 1610612752 1610612762 1610612751 1610612756 1610612746 1
610612749 1610612743 1610612755 1610612756 ...
## $ VISITOR_TEAM_ID : int 1610612764 1610612737 1610612763 1610612738 1610612747 1610612742 1
610612748 1610612757 1610612764 1610612747 ...
## $ SEASON : int 2020 2020 2020 2020 2020 2020 2020 2020 2020 2020 ...
## $ TEAM_ID_home : int 1610612755 1610612752 1610612762 1610612751 1610612756 1610612746 1
610612749 1610612743 1610612755 1610612756 ...
## $ PTS_home : int 120 101 141 130 102 121 132 128 125 99 ...
## $ FG_PCT_home : num 0.557 0.383 0.544 0.523 0.465 0.536 0.489 0.535 0.495 0.465 ...
## $ FT_PCT_home : num 0.684 0.739 0.774 0.955 0.933 0.9 0.9 0.8 0.697 0.833 ...
## $ FG3_PCT_home : num 0.429 0.364 0.487 0.447 0.308 0.394 0.415 0.429 0.313 0.321 ...
## $ AST_home : int 26 15 28 31 21 23 34 29 27 24 ...
## $ REB_home : int 45 54 42 46 31 39 61 35 40 47 ...
## $ TEAM_ID_away : int 1610612764 1610612737 1610612763 1610612738 1610612747 1610612742 1
610612748 1610612757 1610612764 1610612747 ...
## $ PTS_away : int 95 92 129 108 109 127 98 109 118 90 ...
## $ FG_PCT_away : num 0.402 0.369 0.541 0.424 0.45 0.585 0.402 0.479 0.557 0.434 ...
## $ FT_PCT_away : num 0.633 0.818 0.763 0.783 0.871 0.542 0.686 0.821 0.8 0.607 ...
## $ FG3_PCT_away : num 0.091 0.273 0.348 0.353 0.303 0.529 0.286 0.485 0.4 0.269 ...
## $ AST_away : int 22 17 20 23 24 25 20 15 26 19 ...
## $ REB_away : int 40 41 33 43 39 34 36 40 41 33 ...
## $ HOME_TEAM_WINS : int 1 1 1 1 0 0 1 1 1 1 ...
```

```
# calculating mean on the PTS_home column  
# Notice that we get NA for the answer  
# This means we must have missing values in this column  
mean(df0$PTS_home)
```

```
## [1] NA
```

```
# I will remove the na's from the columns that will be used in the model during the data cleaning portion  
mean(df0$PTS_home, na.rm=TRUE)
```

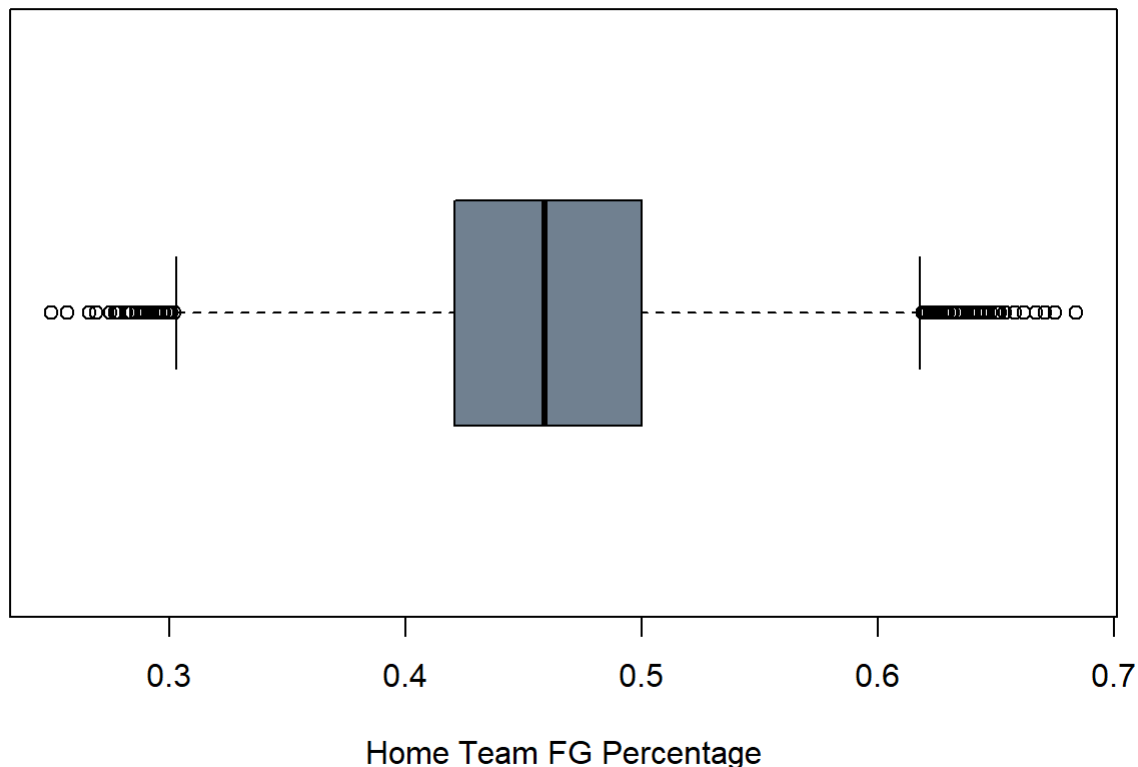
```
## [1] 102.7666
```

```
# Get mean of PTS scored for away teams  
mean(df0$PTS_away, na.rm=TRUE)
```

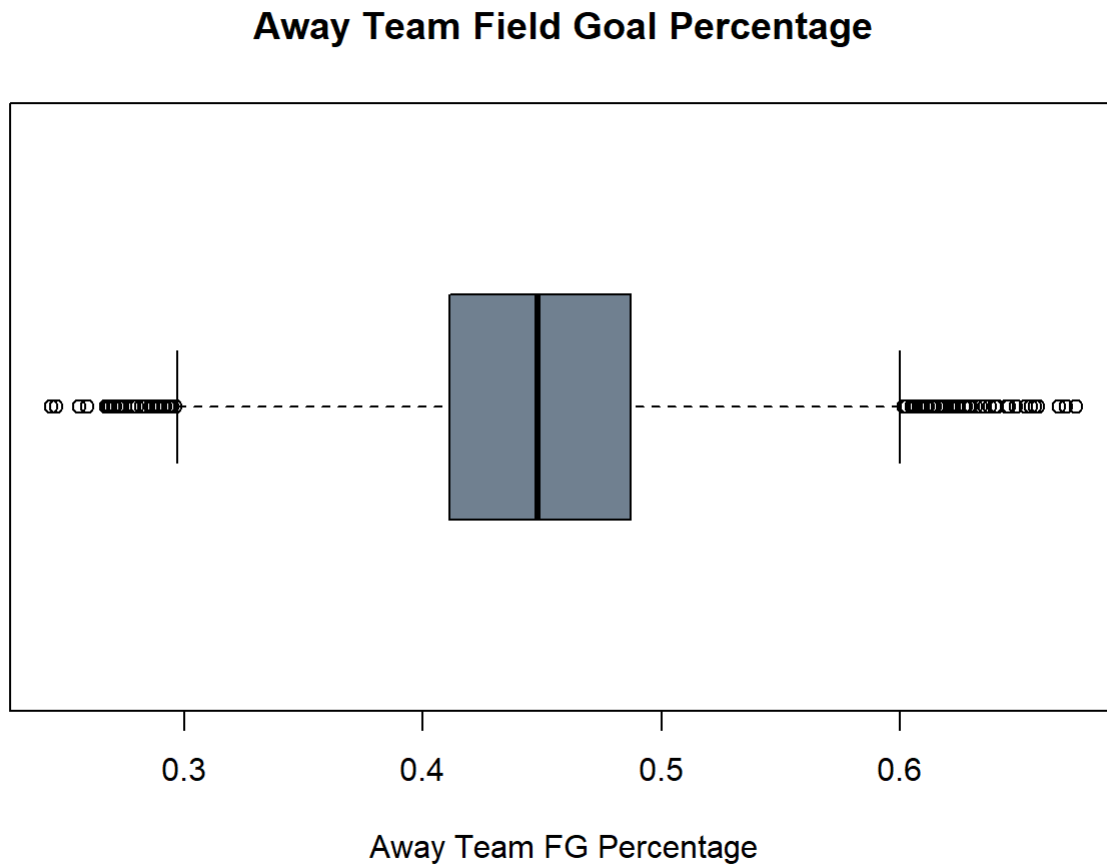
```
## [1] 99.90764
```

```
boxplot(df0$FG_PCT_home, col="slategray", horizontal=TRUE, xlab="Home Team FG Percentage",  
main="Home Team Field Goal Percentage")
```

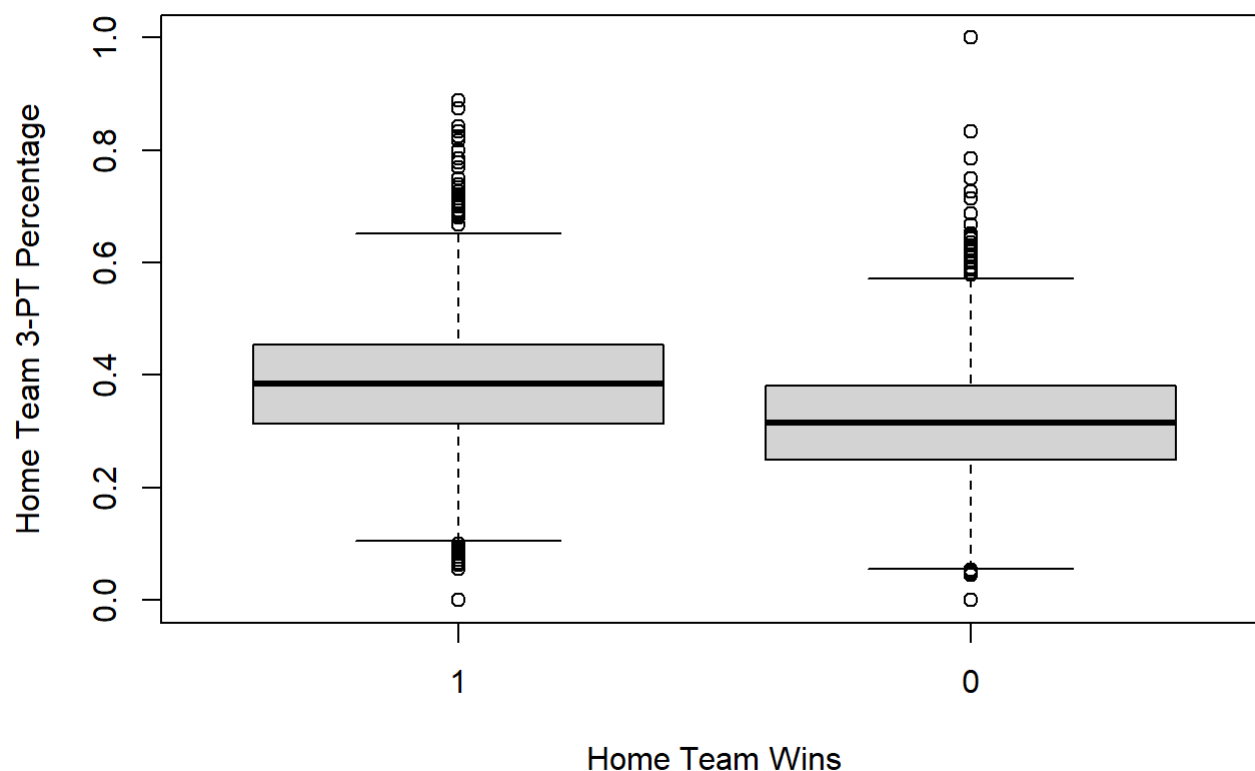
Home Team Field Goal Percentage



```
boxplot(df0$FG_PCT_away, col="slategray", horizontal=TRUE, xlab="Away Team FG Percentage",
main="Away Team Field Goal Percentage")
```



```
# making column factor so we can plot
df0$HOME_TEAM_WINS <- factor(df0$HOME_TEAM_WINS, levels=c("1", "0"))
plot(df0$HOME_TEAM_WINS, df0$FG3_PCT_home, xlab="Home Team Wins", ylab="Home Team 3-PT Percentage")
```



*#There seems to be a relationship between the home team winning and the home team's
3 point percentage*

Data Cleaning

- Link to data: <https://www.kaggle.com/nathanlauga/nba-games> (<https://www.kaggle.com/nathanlauga/nba-games>)
- describe what steps you had to do for data cleaning (more points for messier data that needed cleaning)

I Performed the following steps: 1) Deleted columns that were not needed or unable to help us predict winning teams 2) count the NAs in each of the columns 3) Replace NA with mean

```
# The following columns are not very useful for what we are trying to accomplish
# which is, predicting who will win
```

```
#game_id: col2
#Game_Status_Text : col 3
#home_team_id id : col 4
#visitor_team_id: col 5
#team_id_home : 7
#team_id_away: 14
#game status

#count NAs
sapply(df0, function(x) sum(is.na(x)))
```

```
##    GAME_DATE_EST      GAME_ID GAME_STATUS_TEXT    HOME_TEAM_ID
##          0          0          0          0
## VISITOR_TEAM_ID      SEASON    TEAM_ID_home    PTS_home
##          0          0          0          99
##    FG_PCT_home    FT_PCT_home    FG3_PCT_home    AST_home
##          99          99          99          99
##    REB_home    TEAM_ID_away    PTS_away    FG_PCT_away
##          99          0          99          99
##    FT_PCT_away    FG3_PCT_away    AST_away    REB_away
##          99          99          99          99
##    HOME_TEAM_WINS
##          0
```

```
# remove columns that are unnecessary
df <- df0[-c(2,3,4,5,7,14)]
```

```
# removing those columns coincidentally made it easier to replace NAs in the columns that contain
ed NAs
```

```
df$PTS_home[is.na(df$PTS_home)] <- mean(df$PTS_home, na.rm=TRUE)
df$FG_PCT_home[is.na(df$FG_PCT_home)] <- mean(df$FG_PCT_home, na.rm=TRUE)
df$FT_PCT_home[is.na(df$FT_PCT_home)] <- mean(df$FT_PCT_home, na.rm=TRUE)
df$FG3_PCT_home[is.na(df$FG3_PCT_home)] <- mean(df$FG3_PCT_home, na.rm=TRUE)
df$AST_home[is.na(df$AST_home)] <- mean(df$AST_home, na.rm=TRUE)
df$REB_home[is.na(df$REB_home)] <- mean(df$REB_home, na.rm=TRUE)
df$PTS_away[is.na(df$PTS_away)] <- mean(df$PTS_away, na.rm=TRUE)
df$FG_PCT_away[is.na(df$FG_PCT_away)] <- mean(df$FG_PCT_away, na.rm=TRUE)
df$FT_PCT_away[is.na(df$FT_PCT_away)] <- mean(df$FT_PCT_away, na.rm=TRUE)
df$FG3_PCT_away[is.na(df$FG3_PCT_away)] <- mean(df$FG3_PCT_away, na.rm=TRUE)
df$AST_away[is.na(df$AST_away)] <- mean(df$AST_away, na.rm=TRUE)
df$REB_away[is.na(df$REB_away)] <- mean(df$REB_away, na.rm=TRUE)
```

```
#show na's are deleted
sapply(df, function(x) sum(is.na(x)))
```

```
##  GAME_DATE_EST      SEASON      PTS_home      FG_PCT_home      FT_PCT_home
##              0              0              0              0              0
##  FG3_PCT_home      AST_home      REB_home      PTS_away      FG_PCT_away
##              0              0              0              0              0
##  FT_PCT_away      FG3_PCT_away      AST_away      REB_away      HOME_TEAM_WINS
##              0              0              0              0              0
```

Divide train/test

- Divide into 75/25 train/test, using seed 1234

```
# your code here

set.seed(1234)
i <- sample(1:nrow(df), .75*nrow(df), replace=FALSE)
train <- df[i,]
test <- df[-i,]
```

Algorithm 1: Logistic Regression

- code to run the algorithms
- commentary on feature selection you selected and why
- code to compute your metrics for evaluation as well as commentary discussing the

Commentary on Features Chosen: I decided to predict based on the field goal percentage of both 2-pointers and 3 pointers for each team. Winning comes down to who scores more points. However, predicting simply on the number of points is too simple. The number of points scored can vary due to the quality of defence from team to team. The percentage of made shots will vary less.

Commentary on Results: The algorithm was able to predict the home team winning with only 50% accuracy.

```
attach(df)
glm1 <- glm(HOME_TEAM_WINS~FG_PCT_home+FG_PCT_away+FG3_PCT_home+FG3_PCT_away, data=train, family
="binomial")

summary(glm1)
```



```
##
## Call:
## glm(formula = HOME_TEAM_WINS ~ FG_PCT_home + FG_PCT_away + FG3_PCT_home +
##      FG3_PCT_away, family = "binomial", data = train)
##
## Deviance Residuals:
##      Min        1Q    Median        3Q        Max
## -3.1512  -0.6237  -0.2079   0.6176   3.2858
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -0.5941     0.2288  -2.596  0.00942 **
## FG_PCT_home  -25.8019     0.5110 -50.495 < 2e-16 ***
## FG_PCT_away   26.4695     0.5164  51.259 < 2e-16 ***
## FG3_PCT_home  -4.2062     0.2102 -20.010 < 2e-16 ***
## FG3_PCT_away   4.0325     0.2086  19.332 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 25031  on 18506  degrees of freedom
## Residual deviance: 15087  on 18502  degrees of freedom
## AIC: 15097
##
## Number of Fisher Scoring iterations: 5
```

```
probs <- predict(glm1, newdata=test, type="response")
pred <- ifelse(probs>.5,2,1)
log_reg_acc <- mean(pred==test$HOME_TEAM_WINS)
print(paste("accuracy = ", log_reg_acc))
```

```
## [1] "accuracy = 0.50194489465154"
```

```
table(pred,test$HOME_TEAM_WINS)
```

```
##
## pred    1    0
##    1 3097  661
##    2  490 1922
```

```
detach()
```

Alogorithm 2: Naive Bayes

- code to run the algorithms
- commentary on feature selection you selected and why
- code to compute your metrics for evaluation as well as commentary discussing the results

Commentary on Features Chosen: I decided to predict based on the field goal percentage of both 2-pointers and 3 pointers for each team. Winning comes down to who scores more points. However, predicting simply on the number of points is too simple. The number of points scored can vary due to the quality of defence from team to team. The percentage of made shots will vary less.

Commentary of Results: The Naive Bayes model had an accuracy of 80% which is very good. State of the art NBA predicting projects have an accuracy around 85%.

```
attach(df)
set.seed(1234)
i <- sample(1:nrow(df), .75*nrow(df), replace=FALSE)
train2 <- df[i,]
test2 <- df[-i,]

library(e1071)

nb1 <- naiveBayes(HOME_TEAM_WINS~FG_PCT_home+FG_PCT_away+FG3_PCT_home+FG3_PCT_away, data=train2)
nb1
```

```
##
## Naive Bayes Classifier for Discrete Predictors
##
## Call:
## naiveBayes.default(x = X, y = Y, laplace = laplace)
##
## A-priori probabilities:
## Y
##      1      0
## 0.591614 0.408386
##
## Conditional probabilities:
##   FG_PCT_home
## Y      [,1]      [,2]
## 1 0.4802080 0.05260234
## 0 0.4315018 0.04880490
##
##   FG_PCT_away
## Y      [,1]      [,2]
## 1 0.4288209 0.05001222
## 0 0.4779595 0.05057603
##
##   FG3_PCT_home
## Y      [,1]      [,2]
## 1 0.3833086 0.1106196
## 0 0.3166135 0.1033954
##
##   FG3_PCT_away
## Y      [,1]      [,2]
## 1 0.3224652 0.1056632
## 0 0.3876930 0.1075137
```

```
p1 <- predict(nb1, newdata=test2, type="class")
table(p1, test2$HOME_TEAM_WINS)
```

```
##
## p1      1      0
##  1 3089  697
##  0  498 1886
```

```
naive_acc <- mean(p1==test$HOME_TEAM_WINS)
print(naive_acc)
```

```
## [1] 0.8063209
```

```
detach(df)
```

Algorithm 3: SVM

- code to run the algorithms
- commentary on feature selection you selected and why
- code to compute your metrics for evaluation as well as commentary discussing the results

Commentary on Features Chosen: I decided to predict based on the field goal percentage of both 2-pointers and 3 pointers for each team. Winning comes down to who scores more points. However, predicting simply on the number of points is too simple. The number of points scored can vary due to the quality of defence from team to team. The percentage of made shots will vary less.

Commentary of Results: The SVM model had an accuracy of 81% which is the highest out of all 3 models.

```
attach(df)
set.seed(1234)
i <- sample(1:nrow(df), .75*nrow(df), replace=FALSE)
train3 <- df[i,]
test3 <- df[-i,]

library(e1071)
svm1 <- svm(HOME_TEAM_WINS~FG_PCT_home+FG_PCT_away+FG3_PCT_home+FG3_PCT_away, data=train3, kernel="linear", cost=10, scale=TRUE)
detach(df)
```

SVM Results

```
summary(svm1)
```

```
##
## Call:
## svm(formula = HOME_TEAM_WINS ~ FG_PCT_home + FG_PCT_away + FG3_PCT_home +
##      FG3_PCT_away, data = train3, kernel = "linear", cost = 10, scale = TRUE)
##
##
## Parameters:
##      SVM-Type:  C-classification
##      SVM-Kernel: linear
##              cost: 10
##
## Number of Support Vectors: 8367
##
## ( 4183 4184 )
##
##
## Number of Classes: 2
##
## Levels:
## 1 0
```

```
pred3 <- predict(svm1, newdata=test3)
table(pred3,test3$HOME_TEAM_WINS)
```

```
##
## pred3      1      0
##      1 3082  643
##      0  505 1940
```

```
svm_acc <- mean(pred3==test3$HOME_TEAM_WINS)
svm_acc
```

```
## [1] 0.8139384
```

Results analysis

- rank the algorithms from best to worst performing on your data
- add commentary on the performance of the algorithms
- your analysis concerning why the best performing algorithm worked best on that data
- commentary on what your script was able to learn from the data (big picture) and if this is likely to be useful

Ranking Commentary: As we can see from the results SVM had the highest accuracy at 81.4%. Naive Bayes came in at a close second with 80.6% and logistic regression scored the lowest with only 50.2% accuracy.

Commentary on Performance: Logistic Regression: The algorithm was the worst performing out of the three with only 50% accuracy. It correctly classified 3097 positive values and 1922 negative values. While incorrectly classifying 661 positive values and 490 negative values.

Naive Bayes: Naive Bayes had the 2nd highest accuracy at 80.6%. The model was able to correctly classify 3089 positive values and 1886 negative values. While incorrectly classifying 498 negative values and 687 positive values.

SVM: SVM had the highest accuracy score with 81%. The model correctly classified 3082 positive values and 1940 negative values. While misclassifying 505 negative values and 643 positive values.

Analysis on best performing algorithm: The fact that my SVM model was the top performing algorithm was not surprising to me. SVMs are used by the top performing NBA game predicting projects. The reason why SVMs work well with data sets such as NBA statistics is because they contain many interdependent and related features. For example, points scored is somewhat related to the amount of shots taken, or the number of assists is somewhat intertwined with the number of points scored, etc. SVMs are able to create multi-dimensional feature vectors which means they are capable of capturing the interactions between these related features in the statistics.

Big Picture: What the script was able to find from the data is that sports statistics, such as NBA statistics for this data set, have many intertwined statistics that have some relationship with one another. Because of this, some ML algorithms can have a difficult time correctly classifying results. When it comes to dealing with data sets with many related features, such as statistics for sports, Support Vector Machines are a good option to deal with these related features due to their ability to create multi-dimensional feature vectors.

```
print(paste("Logistic Regression accuracy", log_reg_acc))
```

```
## [1] "Logistic Regression accuracy 0.50194489465154"
```

```
print(paste("Naive Bayes accuracy", naive_acc))
```

```
## [1] "Naive Bayes accuracy 0.806320907617504"
```

```
print(paste("Support Vector Machine (SVM) accuracy", svm_acc))
```

```
## [1] "Support Vector Machine (SVM) accuracy 0.813938411669368"
```